Emotion-Aware, Agentic Healthcare Chatbot Proposal

Prepared by: Rishabh Gupta — July 2025

1 Problem Statement

Modern tele-health interactions lack two critical capabilities:

1) **Emotional Intelligence:** Traditional chatbots "hear" what patients say but ignore how they say it—anxiety and distress go unnoticed.

2) **Modular Expertise & Memory:** Monolithic dialogue models struggle with medical accuracy and conversational empathy, and lack persistent memory across sessions.

Goal: Build a real-time, multi-agent healthcare chatbot that:

- Detects and adapts to patient emotion in real time.
- Maintains a multimodal memory of content and affective state.
- Extracts symptoms, performs RAG-powered medical lookups, suggests diagnoses.
- Recommends and books appointments with nearby doctors based on location and specialty.

2 Solution Overview

Our design combines **LangGraph** orchestration, specialized LLM agents (OpenAI/Anthropic), speech modules (VAD, STT, TTS), and a high-performance vector memory store.

1. **Emotion Detection Agent**: Monitors voice cues (tone, pitch, pauses) via STT and paralinguistic analysis; tags each snippet with valence/arousal labels.
2. **Multimodal Memory Agent**: Streams transcripts and emotion metadata into a vector DB (e.g. FAISS); supports short-term and long-term recall.
3. **Specialized LLM Agents**:
   - Symptoms Extraction Agent
   - Medical Retrieval Agent (RAG over curated medical knowledge)
   - Empathy Response Agent (modulates tone based on emotional state)
   - Orchestrator Agent (task routing, parallel execution, error handling)
4. **Output & Action**: Presents top-3 differential diagnoses and suggests/book appointments with specialists via integrated scheduling APIs.

3 System Architecture

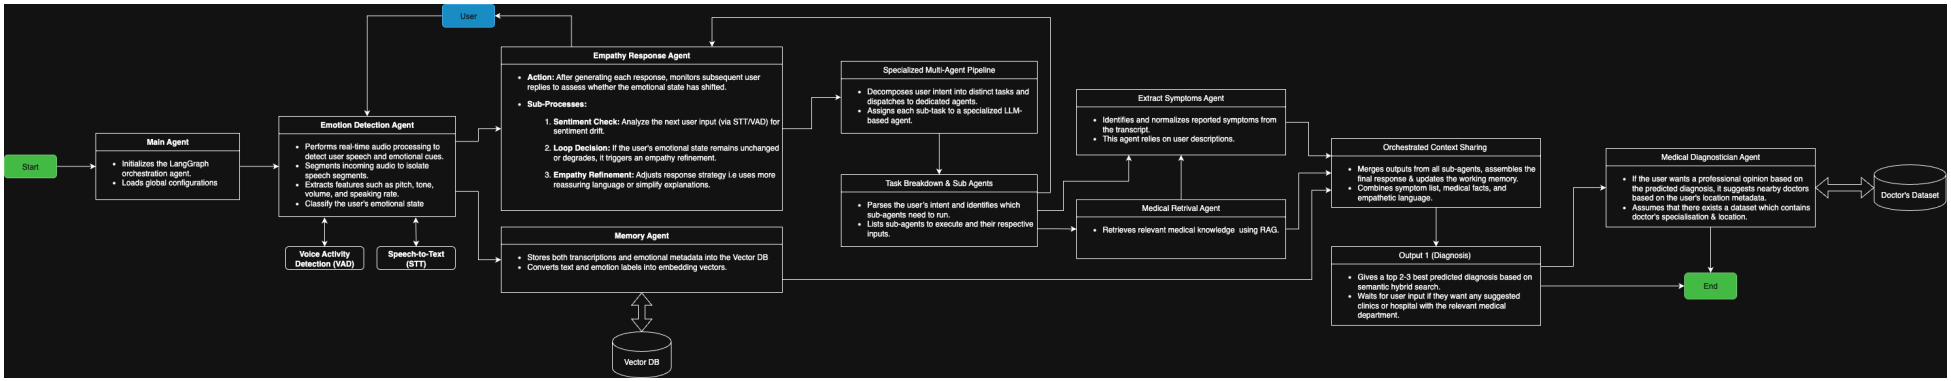

Figure 1: Agentic Chatbot System Architecture (Rotated)

3.1 High-Level Flow

1. User speaks or types input.
2. Emotion Detection Agent: VAD → STT → paralinguistic analysis.
3. Memory Agent: Ingests transcript + emotion tags into vector DB.
4. LangGraph Orchestration: Decomposes request, dispatches sub-agents in parallel.
5. Symptoms Extraction → Medical Retrieval (RAG lookup).
6. Empathy Response Generation with iterative sentiment loop.
7. Diagnosis presentation + doctor recommendation + booking.
8. Feedback loop: Next user utterance rescored for sentiment drift.

3.2 Key Modules

Module	Responsibility
STT & VAD	Convert speech to text; detect voice activity, pitch, tone, and pauses.
Emotion Detection Agent	Classify valence/arousal; trigger empathy refinement if distress persists.
Memory Agent	Store multimodal embeddings in vector DB; support retrieval-augmented prompts.
LangGraph Orchestrator	Define agent graph; handle task decomposition, parallelism, context sharing, and fallbacks.
Symptoms Extraction Agent	Normalize free-form input into structured symptom lists.
Medical Retrieval Agent	Query curated clinical guidelines and disease ontologies via RAG.
Empathy Response Agent	Generate responses modulated by real-time emotion checks.
Appointment Booking	Suggest specialists (geolocation + doctor dataset); integrate scheduling APIs.

4 Core Features & User Experience

- **Emotion-Adaptive UX:** On-screen emotion meter, seamless shifting between technical and empathetic modes.
- **Persistent Medical History & Emotions:** Recall past diagnoses, medications, and emotional context for tailored follow-ups.
- **Symptom-to-Diagnosis Pipeline:** NL symptom extraction → RAG lookup → top-N differential diagnoses with confidence scores.
- **Doctor Recommendation & Booking:** From symptom report to scheduled appointment in three conversational turns.

5 Technology Stack

Layer	Technology
Orchestration	LangGraph
LLM APIs	OpenAI GPT-4, Anthropic Claude
Speech Modules	WebRTC VAD, Whisper STT, Amazon Polly / Azure TTS
Vector Memory	FAISS or Pinecone
Backend Framework	Python, FastAPI, Docker, Kubernetes
Databases	PostgreSQL, Redis, Vector DB
Frontend	React (Web), Native Mobile SDKs
Monitoring & Logging	Prometheus, Grafana, ELK
Scheduling API	Calendly; Hospital Scheduling Services

6 Scalability, Security & Privacy

- **Horizontal Scaling:** Containerized LangGraph workflows, Kubernetes HPA.
- **Failover & Redundancy:** Multi-region LLM endpoints; fallback to local models.
- **Data Privacy:** HIPAA/GDPR compliance; end-to-end encryption; user controls for memory management.

7 Next Steps & Prototype Plan

1. **MVP (6 weeks):** Core STT → Emotion tagging → LangGraph orchestration → symptom extraction → RAG → simple empathy loop; basic UI.
2. **Pilot Testing (4 weeks):** 20-user trial; validate emotion accuracy (target 80% human agreement); appointment flow usability.
3. **Iteration (8 weeks):** Add parallel multi-model generation, advanced empathy refinements, admin memory dashboard, full scheduling integration.

8 Conclusion

Our proposed architecture—combining multimodal emotion memory, a LangGraph agent pipeline, and RAG-powered medical retrieval—yields an empathetic, accurate, and action-oriented healthcare chatbot. We look forward to prototyping and demonstrating its impact on patient engagement and care efficiency.